Second Brain multi-page parser fixture, page one of three.
This page proves that the parser walks past the first page boundary and records page one content.
Alpha paragraph: the local owner keeps invoices, contracts and retention notes in a private corpus.
The quick brown fox jumps over the lazy dog while the gateway token stays on loopback.
Page one reference code is ALPHA-PAGE-ONE-2049 and it appears nowhere else in this docume
Every line here is ordinary text so the engine extracts it without optical character recognition.

Second Brain multi-page parser fixture, page two of three.
This page proves that the parser continues after the first page and records page two content.
Beta paragraph: contracts renew each quarter and the retention policy keeps active rows only.
A different sentence about warehouse receipts, shipping notes and signed delivery records follows.
Page two reference code is BETA-PAGE-TWO-3310 and it appears nowhere else in this document.
The engine should place these lines on page two with bounding boxes measured from the page.

Second Brain multi-page parser fixture, page three of three.
This page proves that the parser reaches the last page and records page three content.
Gamma paragraph: the final audit trail names a retention window and a review date for the owner.
Closing sentence about quarterly reconciliation, archived receipts and the signed summary letter.
Page three unique marker zebrafinch7731 identifies the final page in retrieval results.
If retrieval returns this marker, every page of the document was extracted, not only the first.